Exam 1 Practice Problems

Part 1 - Classification Algorithms

After the first exam in a data mining course, the results of the exam were recorded along with some information about each student. The data is below:

ID	Passed All Assignments	GPA	Language	Passed Exam
1	No	3.1	Python	Yes
2	No	2.0	Python	No
3	Yes	3.5	C++	Yes
4	Yes	2.5	Java	Yes
5	Yes	3.9	Python	No
6	No	2.9	C++	No
7	Yes	3.2	Java	Yes

1. Using a KNN classifier with K=3, predict whether the following student will pass the exam. (Do not worry about normalizing the data.)

8	Yes	3.0	C++	?

Yes

2. Using a Naive Bayes classifier, predict whether the student will pass the exam. Bin the GPA feature into >=3.0 and <3.0

$P(Yes|Yes) \cdot P(>=3.0|Yes) \cdot P(C++|3.0) \cdot P(Yes)$

$3/4 \quad 3/4 \quad 1/4 \quad 4/7 = \frac{9}{16} \cdot \frac{1}{7} = Yes$

$1/3 \quad 1/3 \quad 1/3 \quad 3/7 = \frac{1}{9} \cdot \frac{1}{7}$

3. Given the following dataset:
Different tissue papers & whether or not they are good for your science experiment.
(Yes, the color matters in this problem.)

ID #	Color	Acid Durability	Strength	Class
1	Yellow	7	7	bad
2	White	7	4	good
3	Yellow	3	4	good
4	Green	1	4	good
5	White	5	5	bad
6	White	6	3	bad

If you want to create a decision tree to classify the data, what is the best attribute to split on first?

- **Use Gini index as the measure of impurity**
- **Also know how to use entropy as the measure of impurity, either one is fair game for the exam!**

Note: This problem is too long for an exam, so I won't ask you to do something this long on the exam. But you do need to know how to do this - it'll just be something shorter on the exam.

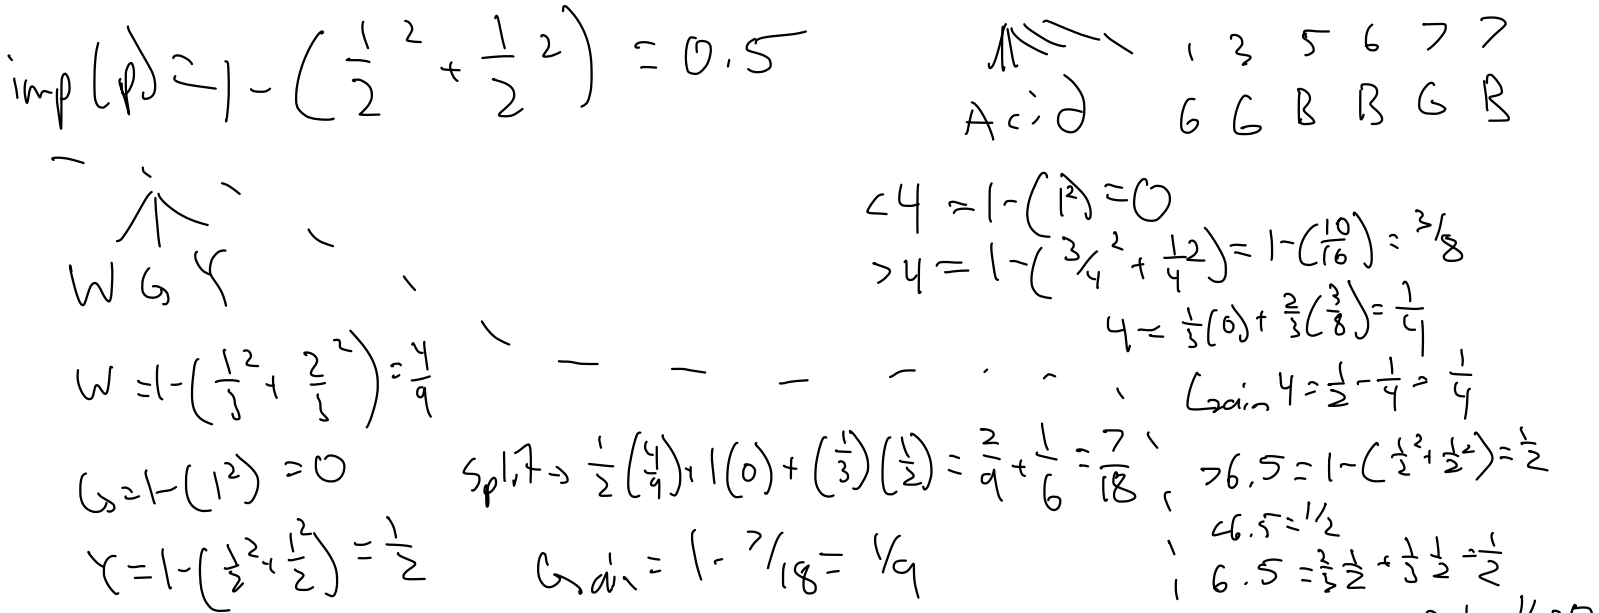

SVMs: Make sure you understand the SVM practice problem questions!

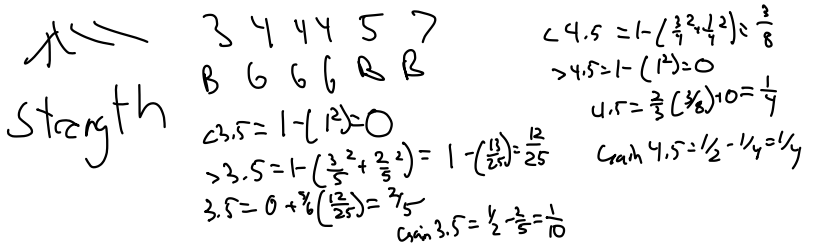

Part 2 - Linear Regression

A scientist is researching whether or not birds exposed to pollutants lay eggs with thinner shells. She collects a sample of egg shells from 5 different nests and measures the pollution level and thinness of the shell. Her results are below:

Pollution	3	8	30	25	15
Thinness	1	3	9	10	5

1. Find the equation of the regression line for this data.

$\bar{Y} = 28/5 \quad \bar{X} = 81/5$

$\beta_1 = 0.330 \quad \beta_0 = 0.259$

2. Calculate the R2 of the line.

$$\frac{var(m) - var(r)}{var(m)} = \frac{11.84 - 0.724}{11.84} = 0.9388$$

3. Calculate the RMSE of the line.

0.8509

Part 3 - Evaluating Classifiers

Given the following confusion matrices for two different classifiers:

Classifier 1		Predicted	
		+	-
Actual	+	50	20
	-	130	300

Classifier 2		Predicted	
		+	-
Actual	+	60	10
	-	30	400

1. Which classifier is better on the basis of error rate?

$\frac{150}{500} = 30\%$ $\quad \frac{40}{500} = 8\%$

✓ ↗

2. Which classifier is better on the basis of F-measure (for the positive class only)?

$P = \frac{50}{180}, \frac{60}{90} \qquad r = \frac{50}{70}, \frac{60}{70}$

$F_1 = \frac{2 \cdot \frac{50}{180} \cdot \frac{50}{70}}{\frac{50}{180} + \frac{50}{70}}$ ↓ 2/5

$F_2 = \frac{2 \cdot \frac{60}{90} \cdot \frac{60}{70}}{\frac{60}{90} + \frac{60}{70}}$ ↙ 3/4

✓ ↗

Part 4 - Short Answers

1. What is the difference between noise and outliers?

↗ ↑ unusual data

random errors and variations

2. Give 2 ways of dealing with missing values in a dataset.

Imputation, Deletion

3. What is the curse of dimensionality?

more dim → more sparse

4. What is overfitting and why is it a problem?

Low training error but high test error

5. What is the naive assumption in Naive Bayes?

independence

6. Describe and/or draw a situation in which using unweighted voting for KNN gives you a different classification than weighted voting.

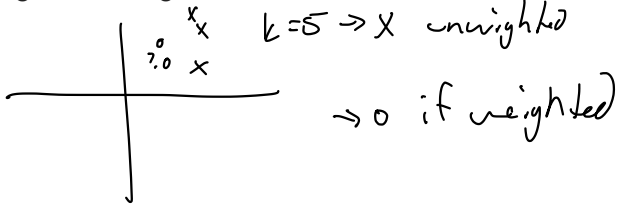

7. Explain "slack" in an SVM - what is it and why do we need slack variables?

To allow misclassifications,

Reduce overfitting